

Does the Conformal Guarantee Travel? Measuring, Auditing and Repairing Coverage Loss in Cross-Dataset Trajectory Prediction

Hassan Waqas

Abstract—[TODO: Write the abstract LAST, from the generated numbers: problem (a conformal guarantee calibrated on one driving dataset is deployed on another with zero labels), what was measured (gap, models, directions), the diagnosis (covariate reweighting fails; identifiability), the solution(s) that actually worked with their measured effect, and the limits. Do not state any effect that is not in generated/numbers.tex.]

Index Terms—Conformal prediction, trajectory prediction, distribution shift, uncertainty quantification, autonomous driving.

AUTONOMOUS vehicles plan around forecasts of what nearby road users will do. A forecast is only as useful as the statement it makes about its own reliability: a planner that knows a predicted region contains the true future with probability at least $1 - \alpha$ can keep clear of that region and of nothing larger. Conformal prediction (CP) supplies such a statement for any black-box predictor, with a finite-sample guarantee that requires nothing beyond exchangeability of calibration and test data [1], [2], and it has been adopted for trajectory prediction and safe planning [3]–[5].

Exchangeability is exactly what deployment breaks. A forecaster is calibrated on the data at hand—one set of cities, one sensor suite, one annotation pipeline—and then driven somewhere else. That accuracy drops when a trajectory predictor moves between datasets is documented [6], and that predictive uncertainty degrades under dataset shift is well known for classifiers [7]. What is not known is what happens to the *guarantee itself* when a conformally calibrated forecaster is moved, without any new labels, to a different real driving dataset: whether the stated $1 - \alpha$ survives, how badly it fails, and whether the standard remedies for shift repair it.

The conformal literature offers several remedies, each resting on an assumption. Weighted CP restores validity under covariate shift when the likelihood ratio is known or can be estimated [8]; more general results bound the coverage loss by a distance between calibration and test distributions [9]; online methods track shift by feeding realised errors back into the threshold [10], [11], an idea carried to trajectory prediction [5] and to robotic settings more broadly [12]; and a generative-prior approach has been proposed for planning under environment shift [13]. What these do not establish is the *static, label-free, dataset-to-dataset* setting in which a forecaster is frozen, calibrated once, and deployed: how large the failure is on real data, whether the covariate-shift remedy

suffices, and how many labelled target scenes it takes to detect or repair the failure.

Contributions: [TODO: Finalise after results. Planned, each to be kept ONLY if the pre-registered test supports it: (i) a measurement protocol and the first measurement of zero-label conformal coverage transfer between real driving datasets; (ii) a diagnosis — covariate reweighting does not repair the loss — with a proposition explaining why label-free repair is limited; (iii) the solution(s) that worked: model-uncertainty-normalised scores (H4), a label-efficient coverage audit with a closed-form label budget (H5b), a label-free shift monitor (H6); (iv) released code, pre-registration and experiment log.]

I. RELATED WORK

a) *Conformal prediction for trajectories.*: Split CP turns any predictor into one with marginal coverage guarantees [1], [2] and underlies conformal safe planning around predicted agents [3]. Multi-step and multi-modal trajectories have been handled with copula-based intervals [4], and scenario-aware intervals in a Frenet frame have been proposed on nuPlan [14] (a preprint evaluated on a single dataset). Locally weighted scores that divide the residual by an estimated spread [15] are the classical route to adaptive regions; our normalised score uses the forecaster’s own predicted scale in that role.

b) *Conformal prediction under shift.*: Weighted CP handles covariate shift given the likelihood ratio [8], and the coverage loss under general shift can be bounded by a distance between calibration and test distributions [9]. Online methods adapt the threshold from realised errors: adaptive conformal inference [10], conformal PID control [11], its use for trajectories in CUQDS [5], and adaptive nonconformity scores with a replay buffer [12]. These need a stream of ground-truth labels after deployment. Other work conditions on an observable nuisance parameter and generates synthetic data with a diffusion model [13] or treats the shift the ego vehicle itself induces in other agents [16] (both preprints). We study a different point in this space: no feedback, no generative model, and two real driving datasets.

c) *Cross-dataset trajectory prediction.*: UniTraj harmonises several driving datasets and reports that accuracy drops when models are transferred between them [6]; that study is about point accuracy, not about the validity of an uncertainty guarantee. Beyond trajectories, [7] showed that calibration of deep models degrades under dataset shift; our

question is the conformal analogue for structured, multi-modal forecasts.

II. PROBLEM SETTING

A. Conformal regions for multi-modal trajectories

A forecaster f maps a scene x to K candidate futures $\hat{y}_1, \dots, \hat{y}_K \in \mathbb{R}^{T \times 2}$ for a target agent whose true future is y . We use the nonconformity score

$$s(x, y) = \min_{k \leq K} \max_{t \leq T} \|\hat{y}_{k,t}(x) - y_t\|_2, \quad (1)$$

so that $s \leq q$ iff y lies inside the union of the K tube-shaped balls of radius q around the candidates. Given n calibration scenes from a *source* dataset, split CP sets $\hat{q} =$ the $\lceil (n+1)(1-\alpha) \rceil$ -th smallest score, and the region $C_{\hat{q}}(x) = \{y : s(x, y) \leq \hat{q}\}$ has $\Pr[y \in C_{\hat{q}}(x)] \geq 1 - \alpha$ when the test scene is exchangeable with the calibration scenes [1], [2].

B. The transfer problem

Let P_S, P_T be the source and target scene–future distributions. The *coverage gap* of a source-calibrated threshold on the target is

$$\Delta = (1 - \alpha) - \Pr_{P_T}[s(x, y) \leq \hat{q}], \quad (2)$$

positive when the guarantee fails (under-coverage). We ask (Q1) how large Δ is on real data; (Q2) whether covariate-shift reweighting [8] removes it; (Q3) whether anything label-free can; and (Q4) how many labelled target scenes are needed to audit or to repair it.

C. What label-free repair can and cannot do

Covariate reweighting is valid when $P_T(y | x) = P_S(y | x)$. If the conditional differs—annotation pipelines, sensor noise, and agent behaviour all differ across datasets—no reweighting on x repairs the gap. More strongly, without restricting the class of shifts nothing label-free can.

Proposition 1 (Non-identifiability of a label-free threshold). *Let S_0 be a nonconformity score with a continuous distribution on the target, and consider the scale-shift family $S = c S_0$, $c \geq 1$, whose members are all indistinguishable from unlabeled target inputs. Let $\hat{q}$ be any threshold computed without target labels, so its law does not depend on c . If $\Pr[c S_0 \leq \hat{q}] \geq 1 - \alpha$ for every $c \geq 1$, then $\Pr[\hat{q} = \infty] \geq 1 - \alpha$.*

Proof. For fixed $\hat{q} < \infty$, $\mathbf{1}\{c S_0 \leq \hat{q}\} \rightarrow 0$ pointwise as $c \rightarrow \infty$. By dominated convergence, $\mathbb{E} \Pr[c S_0 \leq \hat{q} | \hat{q}] \rightarrow \Pr[\hat{q} = \infty]$, and the coverage requirement gives $\Pr[\hat{q} = \infty] \geq 1 - \alpha$. $\square$

Remark 1. *The proposition does not say label-free methods are useless; it says any guarantee for all scale shifts is vacuous, so a useful label-free method must either restrict the shift class (as our normalised score does, Sec. III-A) or be complemented by a label-based audit (Sec. III-C).*

D. Cost of using target labels

If k labelled target scenes are used directly, the resulting empirical coverage is Beta-distributed around $1 - \alpha$ [17], with standard deviation about $\sqrt{\alpha(1 - \alpha)/k}$, so ≈ 1000 labels are needed for ± 2 points at $\alpha = 0.1$. Detecting a shortfall δ with an exact one-sided binomial test at level ℓ and power $1 - \beta$ needs

$$k \approx \left(\frac{z_{1-\ell} \sqrt{\alpha(1 - \alpha)} + z_{1-\beta} \sqrt{(\alpha + \delta)(1 - \alpha - \delta)}}{\delta} \right)^2. \quad (3)$$

For $\delta = 0.034$, $\alpha = 0.1$, $\ell = 0.05$, $\beta = 0.1$ this gives $k \approx 750$.

III. METHODS

A. Model-normalised score (label-free)

The forecaster emits, per mode k , a Laplace scale $(b_{k,t}^x, b_{k,t}^y)$. Let $u_k(x) = \frac{1}{T} \sum_t \sqrt{(b_{k,t}^x)^2 + (b_{k,t}^y)^2}$ and use the score $s'(x, y) = \min_k \max_t \|\hat{y}_{k,t} - y_t\| / u_k(x)$. The region is the union of balls of radius $\hat{q}' u_k(x)$. The premise is that when a shift makes scenes harder, the forecaster’s own scale widens, so the normalised score is less shift-sensitive. Its efficiency is measured by region area $\sum_k \pi (\hat{q}' u_k)^2$ at matched coverage.

B. Recalibration with k target labels

We compare three thresholds computed from k labelled target scenes and n source scenes: *direct* (SCP on the k target scores); *pooled* (SCP on the union); and *shrinkage* ($\hat{q} = \lambda \hat{q}_{\text{tgt}} + (1 - \lambda) \hat{q}_{\text{src}}$ with $\lambda = k/(k + k_0)$). We report the label budget k^* at which absolute coverage error falls below 0.02 in at least 90% of random label draws.

C. Coverage audit

Given a source-calibrated $\hat{q}$ and k labelled target scenes, an exact one-sided binomial test of H_0 : coverage $\geq 1 - \alpha$ flags a violation. We report its empirical power and its false-alarm rate on a shift-free split, alongside the theoretical budget of Eq. (3).

D. Label-free shift monitor

A domain classifier on label-free scene features (native sampling interval, nearby-agent count, map density, nearby lanes, history speed, heading change, curvature) is trained to separate source and target scenes; its AUC is the monitor. Within a dataset, the same classifier is trained per city against the rest. We test whether the AUC predicts the coverage loss (across cities, by Spearman correlation).

IV. EXPERIMENTAL SETUP

a) *Data.*: Argoverse 2 [18] and nuScenes [19], converted with ScenarioNet [20] and loaded through UniTraj [6]. Each dataset’s held-out split is divided by a fixed salted hash into a calibration half and a test half, so calibration and test scenes never overlap. The AV2 test split has no labels and is not used.

b) *Forecasters.*: AutoBot-Ego [21] (1.5M parameters), $K = 6$ modes, all trained on CPU (a 5 GB GPU shared with other work was not consistently available); all models are small relative to published state of the art and we state this as a limitation. The forward-direction (AV2-source) model zoo used for the H8 robustness check so far: **av2_cpu_v1**, trained from scratch on 60k AV2 scenes for 10 epochs (selected checkpoint: epoch 8, val minADE₆ 1.349); **av2_valsplit_v1**, trained from scratch on a smaller, disjoint 13.7k-scene split for 20 epochs (epoch 17, val minADE₆ 1.347); **av2_cpu_v2**, fine-tuned from av2_cpu_v1’s checkpoint with a decaying learning rate for 4 further epochs (epoch 3, val minADE₆ 1.092). **[TODO: Add av2_cpu_v2 to the model-zoo results table once results/run2/av2_cpu_v2_forward.json exists (prediction in progress); add the reverse-direction model ns_cpu_v1 (nuScenes-source, architecture-matched) once trained.]**

c) *Protocol.*: $\alpha = 0.1$. Forward direction: calibrate on AV2 calibration half, test on AV2 test half (same-domain control) and on nuScenes (cross-dataset). Reverse direction: the swap. Confidence intervals are percentile bootstraps over scenes. Hypotheses H1–H9 and their pass thresholds were written before the corresponding runs (notes/falsification.md); those defined after seeing Run 1 are marked post-hoc there and are confirmed only on Run 2 data.

d) *Metrics.*: Coverage, coverage gap Δ , region area proxy, label budget k^* , audit power/false-alarm, monitor AUC and Spearman correlation.

V. RESULTS

Status: 2/3 forward-direction models analysed, 0/1 reverse-direction models. Table II and the numbers below are regenerated automatically from results/run2/.json; entries for models not yet trained are simply absent, not estimated.*

A. H1: the coverage gap is real

Calibrated on Argoverse 2 at $\alpha = 0.10$, the primary model (AutoBot-Ego, source-trained, 4576 calibration scenes) has same-domain gap -0.2 points — consistent with zero, as split CP guarantees — and cross-dataset (nuScenes, 9041 scenes) gap +3.3 points (95% CI [2.1, 4.5]). This replicates the preliminary Run 1 measurement (+3.4pt) on an independent calibration/test split. Table II shows a second, more heavily trained model with a substantially *larger* gap (+8.0pt) in the same direction: the failure is reproducible, but its magnitude is not a fixed property of the dataset pair — it depends on the source model. We report this without averaging it away.

B. H2/H3: covariate reweighting does not repair it

Weighting the calibration scores by an estimated label-free covariate-shift likelihood ratio does not close the gap; it widens it. On the primary model the label-free-score gap goes from +3.2 points uncorrected to +5.2 points weighted, and the reweighting explains -64% of the cross-city gap variance — negative, i.e. the weighted fit is worse than assuming no shift

TABLE I
DOES COVARIATE-SHIFT REWEIGHTING REPAIR THE GAP? “UNCORR.” / “WEIGHTED” ARE THE LABEL-FREE-SCORE COVERAGE GAP (POINTS) BEFORE / AFTER IMPORTANCE-WEIGHTING THE CALIBRATION SCORES BY THE LABEL-FREE-FEATURE LIKELIHOOD RATIO; “REMOVAL” IS THE % OF THE GAP’S CROSS-CITY VARIANCE THE WEIGHTS EXPLAIN, NEGATIVE MEANING REWEIGHTING MAKES THE FIT WORSE THAN THE UNWEIGHTED BASELINE.

Model	Dir.	uncorr.	weighted	removal (%)
av2_cpu_v1	forward	3.2	5.2	-64
av2_valsplit_v1	forward	7.9	10.0	-27




Fig. 1. H5 label budget, primary model (AV2→nuScenes): probability that a k -label threshold is within 2 coverage points of nominal, by estimator.

at all (Table I). This is consistent across both models analysed so far and matches the Run 1 finding: it is evidence for conditional, not purely covariate, shift, and is the empirical motivation for Proposition 1 (no label-free method can be expected to repair an unrestricted shift).

C. H4: the normalised score (label-free)

The model-uncertainty-normalised score (Sec. III-A) reduces the primary model’s gap by 49% (from +3.3 to +1.7 points). Table II’s “red.” column shows this is not a uniform effect: on the second model the same normalisation reduces the gap by only 12%. H4 is therefore *not* confirmed as a general repair — it helps substantially on one model and barely on another, and we report both rather than the more favourable one.

D. H5/H5b: recalibration and the coverage audit

At $k = 1000$ labelled target scenes, the exact binomial audit (Sec. III-C) detects the true violation with probability 0.95 on the primary model, against a false-alarm rate of 2.4% on a shift-free control split — matching the pre-registered budget of Eq. (3) ($k \approx 750$ for the analogous $\delta = 0.034$ case) to the same order of magnitude. Figure 1 shows direct recalibration on the primary model reaches 90% probability of being within 2 coverage points of nominal at $k^* \approx 1000$ labelled target scenes; pooling target labels with the (mismatched) source calibration set is markedly less efficient, consistent with H2/H3’s conclusion that the shift is not a simple reweighting.

E. H6/H7: label-free monitor and per-city coverage

Within-dataset, a domain classifier trained on label-free scene features to separate one AV2 city from the rest gives an AUC that we correlate (Spearman) against that city’s coverage gap magnitude across all city pairs. On the primary model $\rho = 0.37$ ($p = 0.04$); on the second model $\rho = 0.10$ ($p = 0.61$, not significant). The pre-registered threshold was $\rho \geq 0.5$: **H6 is not confirmed** on either model. We report this as a negative result — the label-free monitor tracks shift only weakly, and not reliably enough to substitute for the labelled audit of H5b.

F. H8: robustness across a model zoo

Table II is the model-zoo comparison; **[TODO: expand once av2_cpu_v2 (retrained/fine-tuned model) is analysed — currently 2 of 3 models]**. The gap direction (under-coverage) is consistent so far; magnitude is not, which is itself the H8 finding.

G. H9: reverse direction (nuScenes $\rightarrow$ Argoverse 2)

[TODO: Pending: nuScenes-source model training is queued behind the train-split conversion (in progress). Fill in once results/run2/ns_cpu_v1_reverse.json exists.]

VI. DISCUSSION AND LIMITATIONS

A planner that trusts a source-calibrated region on unseen deployment data is, on this evidence, trusting a region that is too small: at $\alpha = 0.10$ the true miscoverage is $\alpha + 0.03$ to $\alpha + 0.08$ depending on the source model, not α . That gap is silent — nothing in the deployed system signals it — which is precisely why we treat measurement and audit as first-class contributions and not a prelude to the “real” fix. Three results bear directly on practice. First, covariate reweighting is not a safe default here: it made the gap worse on both models tested (Sec. V-B), so a practitioner who reweights and stops is worse off than one who does nothing. Second, the label-free normalised score (H4) is not a substitute for labels either — it repaired about half the gap on one model and about a tenth on another, and nothing observed at deployment time (without target labels) currently tells you which regime you are in. Third, the coverage audit (H5b) is the one method that behaved as pre-registered on both models: with on the order of 10^3 labelled target scenes it reliably detects a violation of this size at a low false-alarm rate. The practical recommendation this supports is modest but concrete: budget for a few hundred to one thousand labelled target-domain scenes before trusting a transferred conformal region, and treat covariate reweighting and the normalised score as diagnostics to try, not guarantees to rely on.

a) *Why the effect size varies by model.*: The second model (Table II) was trained on more data with a different schedule and has a substantially larger gap and a normalised score that helps far less. We do not have a causal explanation for this from two data points; it is exactly the kind of claim H8 (a larger model zoo) is designed to either firm up or refute, and we report it as an open question rather than adjudicate it on $n = 2$.

b) *Limitations (fixed).*: (i) Two datasets and one architecture family; (ii) small CPU-trained models, so absolute accuracy is below published state of the art; (iii) the observed gap is an association—its cause (annotation, sampling, agent mix) is not isolated, only a sampling-rate confound was tested; (iv) marginal, not conditional, coverage; (v) hypotheses added after Run 1 are labelled post-hoc; (vi) Proposition 1 covers scale shifts only.

REPRODUCIBILITY

All code, configuration files, pre-registered hypotheses, and the append-only experiment log are in the project repository (paper01-coverage-transfer). Every number and table in this paper is regenerated from the raw result files by `paper/make_results.py`; no result is typed by hand.

REFERENCES

- [1] V. Vovk, A. Gammerman, and G. Shafer, *Algorithmic Learning in a Random World*. Springer, 2005.
- [2] A. N. Angelopoulos and S. Bates, “Conformal prediction: A gentle introduction,” *Foundations and Trends in Machine Learning*, vol. 16, no. 4, pp. 494–591, 2023.
- [3] L. Lindemann, M. Cleaveland, G. Shim, and G. J. Pappas, “Safe planning in dynamic environments using conformal prediction,” *IEEE Robotics and Automation Letters*, vol. 8, no. 8, pp. 5116–5123, 2023.
- [4] S. H. Sun and R. Yu, “Copula conformal prediction for multi-step time series prediction,” in *International Conference on Learning Representations (ICLR)*, 2024.
- [5] H. Huang, S. He, and F. Miao, “CUQDS: Conformal uncertainty quantification under distribution shift for trajectory prediction,” in *Proceedings of the AAAI Conference on Artificial Intelligence*, 2025.
- [6] L. Feng, M. Bahari, K. M. B. Amor, É. Zablocki, M. Cord, and A. Alahi, “UniTraj: A unified framework for scalable vehicle trajectory prediction,” in *European Conference on Computer Vision (ECCV)*, 2024.
- [7] Y. Ovadia, E. Fertig, J. Ren, Z. Nado, D. Sculley, S. Nowozin, J. V. Dillon, B. Lakshminarayanan, and J. Snoek, “Can you trust your model’s uncertainty? Evaluating predictive uncertainty under dataset shift,” in *Advances in Neural Information Processing Systems (NeurIPS)*, 2019.
- [8] R. J. Tibshirani, R. F. Barber, E. J. Candès, and A. Ramdas, “Conformal prediction under covariate shift,” in *Advances in Neural Information Processing Systems (NeurIPS)*, 2019.
- [9] R. F. Barber, E. J. Candès, A. Ramdas, and R. J. Tibshirani, “Conformal prediction beyond exchangeability,” *The Annals of Statistics*, vol. 51, no. 2, pp. 816–845, 2023.
- [10] I. Gibbs and E. J. Candès, “Adaptive conformal inference under distribution shift,” in *Advances in Neural Information Processing Systems (NeurIPS)*, 2021.
- [11] A. N. Angelopoulos, E. J. Candès, and R. J. Tibshirani, “Conformal PID control for time series prediction,” in *Advances in Neural Information Processing Systems (NeurIPS)*, 2023.
- [12] R. Tumu, A. Singh, and R. Mangharam, “AdaptNC: Adaptive non-conformity scores for conformal prediction under distribution shift,” arXiv:2602.01629, 2026.
- [13] K. Rahaman, J. V. Deshmukh, A. R. Hota, and L. Lindemann, “When environments shift: Safe planning with generative priors and robust conformal prediction,” arXiv:2602.12616, 2026.
- [14] Y. Shu, J. Xu, L. Kong, F. Zhang, G. Yin, and C. Sun, “Scenario-aware uncertainty quantification for trajectory prediction with statistical guarantees,” arXiv:2512.05682, 2025.
- [15] J. Lei, M. G’Sell, A. Rinaldo, R. J. Tibshirani, and L. Wasserman, “Distribution-free predictive inference for regression,” *Journal of the American Statistical Association*, vol. 113, no. 523, pp. 1094–1111, 2018.
- [16] A. E. Binny and A. Dixit, “Who moved my distribution? Conformal prediction for interactive multi-agent systems,” arXiv:2511.11567, 2025.
- [17] V. Vovk, “Conditional validity of inductive conformal predictors,” in *Proceedings of the Asian Conference on Machine Learning (ACML)*, ser. Proceedings of Machine Learning Research, vol. 25, 2012, pp. 475–490.

TABLE II

COVERAGE TRANSFER BY MODEL AND DIRECTION, $\alpha = 0.10$. “RED.” = % REDUCTION IN GAP FROM THE LABEL-FREE NORMALISED SCORE (H4); “POWER/FA” = AUDIT REJECT PROBABILITY AT $k = 1000$ TARGET LABELS ON THE SHIFTED / SAME-DOMAIN SPLIT (H5B); ρ = SPEARMAN CORRELATION BETWEEN THE LABEL-FREE DOMAIN-MONITOR AUC AND THE PER-CITY MISCOVERAGE (H6).

Model	Dir.	n_{cal}	n_{tgt}	gap (pt)	95% CI	norm. gap	red. (%)	power/FA @1000	ρ
av2_cpu_v1	forward	4576	9041	3.3	[2.1, 4.5]	1.7	49	0.95/2.4	0.37
av2_valsplit_v1	forward	4576	9041	8.0	[6.4, 9.4]	7.0	12	1.00/1.0	0.10

- [18] B. Wilson, W. Qi, T. Agarwal, J. Lambert, J. Singh, S. Khandelwal, B. Pan, R. Kumar, A. Hartnett, J. K. Pontes, D. Ramanan, P. Carr, and J. Hays, “Argoverse 2: Next generation datasets for self-driving perception and forecasting,” in *NeurIPS Datasets and Benchmarks Track*, 2021.
- [19] H. Caesar, V. Bankiti, A. H. Lang, S. Vora, V. E. Liong, Q. Xu, A. Krishnan, Y. Pan, G. Baldan, and O. Beijbom, “nuScenes: A multimodal dataset for autonomous driving,” in *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2020.
- [20] Q. Li, Z. Peng, L. Feng, Z. Liu, C. Duan, W. Mo, and B. Zhou, “ScenarioNet: Open-source platform for large-scale traffic scenario simulation and modeling,” in *NeurIPS Datasets and Benchmarks Track*, 2023.
- [21] R. Girgis, F. Golemo, F. Codevilla, M. Weiss, J. A. D’Souza, S. E. Kahou, F. Heide, and C. Pal, “Latent variable sequential set transformers for joint multi-agent motion prediction,” in *International Conference on Learning Representations (ICLR)*, 2022.