Does the Conformal Guarantee Travel? Measuring, Auditing and Repairing Coverage Loss in Cross-Dataset Trajectory Prediction

Hassan Waqas Saleem, Rahat Zeeshan Saleem, and Hina Maryam

***Abstract*—A conformal prediction region promises to contain a road user's true future path a fixed fraction of the time. We ask whether that promise survives when a trajectory forecaster is calibrated on one driving dataset and deployed, without new labels, on another. Using forecasters calibrated on Argoverse 2 or nuScenes and evaluated on each other and on Waymo Open Motion, we measure coverage losses of 6 to 25 percentage points at a 90 percent nominal level, on two architectures whose accuracy is within 15 percent of published values; the third dataset was registered in advance. Reweighting by an estimated covariate-shift likelihood ratio makes the loss worse, and we show that no label-free method can repair an unrestricted shift, although the drift of the forecaster's own predicted scale tracks the loss. The loss concentrates in turns, most of all in left-hand traffic. Measured input differences between the datasets reproduce about a quarter of it for one architecture and, with caveats, all of it for the other. Labelled target data can repair it, but labels arrive one scene at a time, and a guarantee that treats agents as independent can fail when scenes contain many agents. A scene-level Learn-then-Test rule with a betting test restores the guarantee at every labelling budget we tested, and returns a region within twice the oracle size from about 150 to 200 labelled scenes. Conformal guarantees for trajectory prediction should therefore not be assumed to transfer across datasets without an audit and a scene-level recalibration.**

***Index Terms*—Conformal prediction, trajectory prediction, distribution shift, autonomous driving.**

AUTONOMOUS vehicles plan around forecasts of what nearby road users will do. A forecast is only as useful as the statement it makes about its own reliability. A planner that knows a predicted region contains the true future with probability at least $1-\alpha$ can keep clear of that region and of nothing larger. Conformal prediction (CP) supplies such a statement for any black-box predictor, with a finite-sample guarantee that needs only exchangeability of calibration and test data [1], [2], and it is now used for trajectory prediction and safe planning [3]–[6].

Deployment breaks exchangeability. A forecaster is calibrated on the data at hand, with one set of cities, one sensor suite and one annotation pipeline, and is then driven somewhere else. Trajectory predictors lose accuracy when moved between datasets [7], [8], and model calibration degrades under dataset shift [9]. What is not known is what happens to the *guarantee itself* when a conformally calibrated forecaster is moved, without new labels, to a different real driving dataset: whether the stated $1-\alpha$ survives, how badly it fails, and whether the standard remedies repair it. Weighted CP needs the likelihood ratio [10], distance-based bounds need the size of the shift [11], [12], and online methods need a stream of labels [13], [14]. None settles the static, label-free, dataset-to-dataset case that Fig. 1 illustrates.

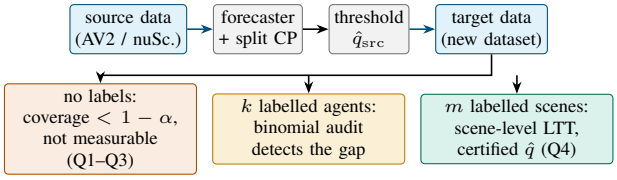


Fig. 1. The transfer problem and the paper's questions. A threshold calibrated on one dataset is deployed on another. Without labels the coverage loss cannot be measured; k labelled agents detect it; m labelled scenes repair it.

Contributions: (i) A direct measurement of zero-label conformal coverage transfer between real driving datasets, in both directions between Argoverse 2 and nuScenes and, registered before it was run, on Waymo Open Motion, which to our knowledge no prior work has quantified at this scale. Coverage falls by 6.4 to 25.0 points at a nominal 90%, confirmed on two architectures within 15% of published accuracy; across training budgets of one architecture the more accurate models lose more coverage, but this does not hold between training seeds or across architectures. (ii) A diagnosis. Covariate reweighting makes the loss worse, a non-identifiability result (Proposition 1) shows why label-free repair is limited in principle, the loss concentrates in turns in left-hand traffic, and controlled input perturbations account for part of it. (iii) A repair protocol, not a new conformal algorithm: it applies Learn-then-Test with a betting test at the scene level, which gives a finite-sample, training-conditional guarantee when labels arrive scene by scene (Propositions 2 and 3). It is valid at every labelling budget we tested, on both architectures, and gives a usable region from about 150 labelled scenes, while the natural agent-level guarantee can fail when scenes contain many agents and the source threshold is already badly wrong. (iv) Code, registered hypotheses and an append-only experiment log, including a defect found and

H. W. Saleem is with a public-sector organization in Pakistan; he received the M.S. degree from the School of Electrical Engineering and Computer Science (SEECS), National University of Sciences and Technology (NUST), Islamabad 44000, Pakistan (corresponding author, e-mail: hsaleem.msee20seecs@seecs.edu.pk; ORCID: 0009-0003-4461-8005).

R. Z. Saleem is with a private-sector organization in Pakistan; he received the M.S. degree from NUST Business School (NBS), National University of Sciences and Technology (NUST), Islamabad 44000, Pakistan (e-mail: rahat.mscm21nbs@student.nust.edu.pk).

H. Maryam is with Kinnaird College for Women University, Lahore, Pakistan (e-mail: hina.maryam@kinnaird.edu.pk).

corrected before submission.

I. Related Work

a) Conformal prediction for trajectories: Split CP turns any predictor into one with marginal coverage guarantees [1], [2] and underlies conformal safe planning around predicted agents [3]. Multi-step and multi-modal trajectories have been handled with copula-based intervals [4], scenario-aware intervals in a Frenet frame on nuPlan [15] and multi-view calibration that fuses vehicle- and infrastructure-side observations [16]; a 2025 survey catalogues these directions and the open gaps around uncertainty and generalisation [17], and a dataset-design study finds that which data the model sees matters more than feature engineering [8], a finding Sec. V-C echoes from the calibration side. Calibrating the predictive *distribution* itself has also been proposed [18]. Locally weighted scores that divide the residual by an estimated spread [19] are the classical route to adaptive regions; our normalised score uses the forecaster's own predicted scale in that role.

b) Safe planning and robot safety with conformal guarantees: Beyond trajectory regions, conformal prediction has certified learned perception for navigation [20], adapted a crowd-navigation policy's safety margin to shift [21], learned safety regions from sparse human feedback [22], certified learned motion primitives [23] and connected learning-enabled control to formal verification [6]. Of these, [24] transfers a safety guarantee across policy updates analytically and [25] treats the shift the ego agent induces in others. All assume that calibration and deployment *domains* match; we study what happens when the domain itself changes.

c) Conformal prediction under shift: Weighted CP handles covariate shift given the likelihood ratio [10], and the coverage loss under general shift can be bounded by a distance between calibration and test distributions [11]. Recent work replaces that bound with an explicit ambiguity-set radius under Lévy–Prokhorov perturbations [12], regularises the representation with a Wasserstein penalty to shrink the achievable coverage gap [26], or uses optimal transport with unlabelled target data to both estimate and mitigate the loss in coverage [27]. That optimal-transport method is the closest label-free competitor, but its bracketing of the unlabelled test scores is built from the minimum and maximum score over all candidate labels, for scores bounded in $[0, 1]$, and is evaluated on classification and a one-dimensional binned regression task. For our unbounded, metre-scale score over multi-modal trajectories that minimum is zero and the maximum is unbounded, and we found no published construction for structured outputs, so we do not run it as a baseline; a trajectory-specific bracketing is a natural extension. Robust conformal prediction instead guards against every test distribution inside an f-divergence ball around the training distribution [28]; it restricts the shift class by construction, as Proposition 1 says a useful label-free method must, but needs a radius chosen or estimated for the shift, and we leave it as a baseline for future work. A parallel line gives sample-complexity lower bounds for when a selective predictor can certify a risk floor under covariate shift at all [29], a negative result in the spirit of our Proposition 1. Online methods adapt the threshold from realised errors: adaptive conformal inference [13], conformal PID control [14], adaptive scores with a replay buffer [30], and variants that revise past predictions [31], play a universal-portfolio game [32] or give training-conditional regret bounds [33]. The closest trajectory-specific instance is CUQDS [5], which pairs a Gaussian-process uncertainty module with online conformal PID control; like the others it needs a stream of ground-truth labels at deployment, so it answers a different question and is not a baseline for our zero-label or fixed-labelled-batch settings. Our ACI comparison (Sec. V-D) covers the same category without CUQDS's architecture-specific module. Other work conditions on an observable nuisance parameter with a diffusion model [34] or treats the shift the ego vehicle induces in others [25] (preprints). We study a different point: no feedback, no generative model, and three real driving datasets.

d) Grouped, hierarchical and conditional coverage: When observations come in groups, exchangeability across observations fails, and hierarchical conformal methods restore validity by pooling or subsampling at the group level [35], [36], with extensions for hierarchical outputs [37], known group-level shift [38] and partial test-group data [39]. The conditional-coverage literature asks for validity on subpopulations rather than on average [40]–[43]. Risk-controlling procedures [44] and Learn-then-Test [45] give finite-sample control of a monotone risk by testing a fixed ordered set of thresholds, and betting-based p-values [46] make such tests far tighter than Hoeffding-type bounds when losses are small. We combine the two for a setting those works do not address: dataset shift in trajectory prediction, where the labelling unit is the *scene* and the controlled risk is the scene-averaged miscoverage of a fixed, source-calibrated forecaster. Valid p-values for bounded losses need care [47]; our own implementation needed a corrected bet cap (Sec. III-D).

e) Risk control and betting, more generally: Risk control has been extended to tradeoffs among several risks [48], semi-supervised tuning [49] and post-processing of decision policies [50]; e-values give anytime-valid alternatives to the p-values we use [51], and weighted martingales monitor post-deployment shift [52]. We use betting for a different purpose: certifying a *scene-averaged* risk from whole labelled scenes.

f) Cross-dataset trajectory prediction and perception under domain shift: UniTraj harmonises several driving datasets and reports that accuracy drops when models are transferred between them [7]; that study is about point accuracy, not about the validity of an uncertainty guarantee. Beyond trajectories, [9] showed that calibration of deep models degrades under dataset shift; our question is the conformal analogue for structured, multi-modal forecasts. The same gap is documented on the perception side: localisation uncertainty is miscalibrated under novel weather [53], collaborative perception generalises poorly across sensor rigs [54], and cross-sensor pretraining is used to close a domain gap that would otherwise need target labels [55], [56]. None of this asks whether a *conformal* guarantee survives the transfer; this paper measures, diagnoses and partially repairs that.

II. Problem Setting

A. Conformal regions for multi-modal trajectories

A forecaster f maps a scene x to K candidate futures $\hat{y}_1, \dots, \hat{y}_K \in \mathbb{R}^{T\times 2}$ for a target agent whose true future is y. We use the nonconformity score

$$s(x, y) = \min_{r\le K} \max_{t\le T} \|\hat{y}_{r,t}(x) - y_t\|_2, \tag{1}$$

so that $s \le q$ iff y lies inside the union of the K tube-shaped balls of radius q around the candidates. Given n calibration scenes from a *source* dataset, split CP sets $\hat{q}$ = the $\lceil (n+1)(1-\alpha)\rceil$-th smallest score, and the region $C_{\hat{q}}(x) = \{y : s(x, y) \le \hat{q}\}$ has $\Pr[y \in C_{\hat{q}}(x)] \ge 1-\alpha$ when the test scene is exchangeable with the calibration scenes [1], [2].

B. The transfer problem

Let P_S, P_T be the source and target scene–future distributions. The *coverage gap* of a source-calibrated threshold on the target is

$$\Delta = (1-\alpha) - \Pr_{P_T}\big[s(x, y) \le \hat{q}\big], \tag{2}$$

positive when the guarantee fails (under-coverage). We ask (Q1) how large Δ is on real data; (Q2) whether covariate-shift reweighting [10] removes it; (Q3) whether anything label-free can; and (Q4) how many labelled target agents are needed to detect it, and how many labelled target scenes to repair it.

C. What label-free repair can and cannot do

Covariate reweighting is valid when $P_T(y \mid x) = P_S(y \mid x)$. If the conditional differs—annotation pipelines, sensor noise, and agent behaviour all differ across datasets—no reweighting on x repairs the gap. More strongly, without restricting the class of shifts nothing label-free can.

Proposition 1 (Non-identifiability of a label-free threshold)**.** *Let S_0 be a nonconformity score with a continuous distribution on the target, and consider the scale-shift family $S = c\,S_0$, $c \ge 1$, whose members are all indistinguishable from unlabelled target inputs. Let $\hat{q}$ be any threshold computed without target labels, so its law does not depend on c. If $\Pr[cS_0 \le \hat{q}] \ge 1-\alpha$ for every $c \ge 1$, then $\Pr[\hat{q} = \infty] \ge 1-\alpha$.*

Proof. For fixed $\hat{q} < \infty$, $\mathbf{1}\{cS_0 \le \hat{q}\} \to 0$ pointwise as $c \to \infty$. By dominated convergence, $\mathbb{E}\,\Pr[cS_0 \le \hat{q} \mid \hat{q}] \to \Pr[\hat{q} = \infty]$, and the coverage requirement gives $\Pr[\hat{q} = \infty] \ge 1-\alpha$. □

Remark 1. *The proposition does not say label-free methods are useless; it says any guarantee* for all *scale shifts is vacuous, so a useful label-free method must either restrict the shift class (as our normalised score does, Sec. III-A) or be complemented by a label-based audit (Sec. III-B).*

D. Cost of using target labels

If k labelled target agents are used directly, the coverage of the recalibrated threshold is Beta-distributed around $1-\alpha$ [57], with standard deviation about $\sqrt{\alpha(1-\alpha)/k}$, so ≈ 900 labels are needed for ± 2 points at $\alpha = 0.1$. Detecting a coverage gap Δ with an exact one-sided binomial test at level a and power $1-\beta$ needs

$$k \approx \left(\frac{z_{1-a}\sqrt{\alpha(1-\alpha)} + z_{1-\beta}\sqrt{(\alpha+\Delta)(1-\alpha-\Delta)}}{\Delta} \right)^2. \tag{3}$$

For $\Delta = 0.034$, $\alpha = 0.1$, $a = 0.05$, $\beta = 0.1$ this gives $k \approx 750$.

III. Methods

A. Model-normalised score (label-free)

The forecaster emits, per mode r, a Laplace scale $(b^x_{r,t}, b^y_{r,t})$. Let $u_r(x) = \frac{1}{T}\sum_t \sqrt{(b^x_{r,t})^2 + (b^y_{r,t})^2}$ and use the score $s'(x, y) = \min_r \max_t \|\hat{y}_{r,t} - y_t\|/u_r(x)$. The region is the union of balls of radius $\hat{q}'\,u_r(x)$. The premise is that when a shift makes scenes harder, the forecaster's own scale widens, so the normalised score is less shift-sensitive. Its efficiency is measured by the area proxy $\sum_r \pi(\hat{q}' u_r)^2$ (the summed cross-section of the tubes, ignoring overlap) at matched coverage; for the unnormalised score this proxy is proportional to $\hat{q}^2$, so the area ratios we report against an oracle equal $(\hat{q}/q_{\text{oracle}})^2$.

B. Coverage audit

Given a source-calibrated $\hat{q}$ and k labelled target agents, an exact one-sided binomial test of H_0: coverage $\ge 1-\alpha$ flags a violation. We report its empirical power and its false-alarm rate on a shift-free split, alongside the theoretical budget of Eq. (3).

C. Certify-or-recalibrate with exchangeable labels

Plain recalibration (split CP on the k labelled target scores alone) is only *marginally* valid: the coverage of the resulting threshold is $\mathrm{Beta}(\ell, k+1-\ell)$-distributed [57], so it falls below $1-\alpha$ in roughly half of all label draws (Sec. V). A practitioner needs the stronger, *training-conditional* statement [57], [58]: coverage $\ge 1-\alpha$ with probability $\ge 1-\delta$ over the draw of the labels. *Certify-or-recalibrate* (C-or-R) delivers it while keeping the source threshold when it can be certified. With k labelled scores $s_1, \dots, s_k$ and $M = \#\{s_i > \hat{q}_{\text{src}}\}$:

1) *Certify.* Keep $\hat{q}_{\text{src}}$ iff the one-sided Clopper–Pearson upper confidence bound [59] on its miss probability at level $\delta/2$ is $\le \alpha$.
2) *Otherwise recalibrate.* Return the ℓ-th smallest labelled score, with ℓ the smallest integer such that $\Pr[\mathrm{Bin}(k, 1-\alpha) \le \ell - 1] \ge 1 - \delta/2$ (the ℓ-th order statistic has coverage $\mathrm{Beta}(\ell, k+1-\ell)$); if no such $\ell \le k$ exists return $+\infty$.

Proposition 2 (C-or-R is training-conditionally valid)**.** *If the k target scores are exchangeable with the target scores at test*

time and continuously distributed, the returned threshold $\hat{q}$ satisfies $\Pr\left[\Pr_{\textit{test}}(s \le \hat{q} \mid \textit{labels}) \ge 1-\alpha\right] \ge 1-\delta$.

Proof. Let F be the target score CDF. The event "$\hat{q}$ fails" is contained in (a) the source threshold is certified although $F(\hat{q}_{\mathrm{src}}) < 1-\alpha$, or (b) the branch-2 order statistic has coverage $< 1-\alpha$. (a) has probability $\le \delta/2$ because the Clopper–Pearson bound is a valid one-sided confidence bound for the miss probability $1-F(\hat{q}_{\mathrm{src}})$ of the fixed threshold; (b) has probability $\le \delta/2$ by the definition of ℓ. A union bound gives δ. □

Certification needs slack: a source threshold calibrated at exactly α has miss probability $\approx \alpha$ and its confidence bound cannot fall below α; in our experiments it is (almost) never certified on shifted data, and on unshifted data it is certified only if calibrated at a smaller $\alpha_{\mathrm{src}} < \alpha$.

D. Scene-level risk control under clustered labels

Driving datasets label many agents per scene, and agents in one scene share road, traffic and annotator. Labels drawn scene by scene are not exchangeable at the agent level, and Proposition 2 then fails (Sec. V). We therefore make the *scene* the unit. For a labelled scene j with agent scores $\{s_{ji}\}$ let $L_j(q) = \frac{1}{n_j}\sum_i \mathbf{1}[s_{ji} > q] \in [0,1]$, non-increasing in q, and control the scene-averaged miscoverage $R(q) = \mathbb{E}_j\, L_j(q)$, which weights every *scene* equally regardless of how many agents it contains. This is a different quantity from the agent-pooled miss rate $\frac{1}{\sum_j n_j}\sum_{j,i}\mathbf{1}[s_{ji} > q]$, which weights every agent equally and so is dominated by whichever scenes happen to be largest; the two coincide only when every scene has the same agent count. We control $R(q)$ because it is the quantity the scene-level test (below) can be made valid for under scene-level exchangeability, and because it is the natural risk when scenes, not agents, are the unit a planner reasons about. Given m labelled scenes, scene-level certify-or-recalibrate proceeds as in Sec. III-C but with Learn-then-Test [45] over a *fixed*, label-independent grid $q_1 > q_2 > \cdots > q_G$ (log-spaced around the source threshold, $\hat{q}_{\mathrm{src}} \times [4, 0.5]$, $G = 300$): for each candidate we compute a p-value for $H_q : R(q) > \alpha$ from the m per-scene losses, test candidates from the largest downwards at level $\delta/2$, and return the last q rejected ($+\infty$ if none). The source threshold is tried first at level $\delta/2$ (certification). Two p-values are used: Hoeffding–Bentkus [44], and a betting p-value [46]: with predictable, variance-adaptive bets $\lambda_i \in [0, \frac{1}{2(1-\alpha)}]$ the capital $K = \prod_i \big(1 + \lambda_i(\alpha - L_i)\big)$ stays non-negative for $L_i \in [0,1]$ and is a supermartingale under $\mathbb{E}L \ge \alpha$, so $p = \min(1, 1/\max_t K_t)$ is valid by Ville's inequality. The cap matters: any bet above $1/(1-\alpha)$ lets a single fully missed scene drive the capital negative, which voids the argument; a regression test in the released code (`test_betting_validity`) checks that the capital never goes negative, that validity holds exactly at the boundary null, and that the old cap is detected.

Proposition 3 (Scene-level guarantee)**.** *If the m labelled scenes are exchangeable draws from the target scene population, and each p-value is valid for its null, the returned threshold satisfies* $\Pr[R(\hat{q}) \le \alpha] \ge 1-\delta$.

Proof. R is non-increasing, so along the descending grid the false nulls ($R(q_i) \le \alpha$) form a prefix and the true nulls a suffix beginning at some i_0. The procedure returns a threshold with $R > \alpha$ only if it rejects $H_{q_{i_0}}$ (fixed-sequence testing stops at the first non-rejection), which has probability $\le \delta/2$ because the grid does not depend on the labels; the certification branch fails with probability $\le \delta/2$. A union bound gives δ. □

Remark 2. *Proposition 3 bounds the scene-averaged miscoverage R, not the agent-weighted miss rate, which can differ when scenes have unequal agent counts; the latter is reported only descriptively (Sec. VI).*

E. Online baseline

Adaptive conformal inference (ACI) [13] updates the working level $\alpha_{t+1} = \alpha_t + \gamma(\alpha - \mathrm{err}_t)$ after every revealed label, keeping the source calibration set fixed; it needs a stream of labels and offers long-run, not finite-sample, coverage. We report the coverage of its threshold frozen after k streamed labels, for $\gamma \in \{0.005, 0.01, 0.05\}$.

F. Controlled input shifts

To locate the source of the loss we perturb one measured AV2–nuScenes difference at a time in AV2 test *inputs* (ground truth untouched) and record the coverage lost by the AV2-calibrated threshold: (i) *2 Hz history*: keep the five 2 Hz keyframes of the 2 s history and interpolate linearly; (ii) *lane dropout*: drop each map polyline with the probability implied by the measured median number of lanes near the target agent (nuScenes 52 vs. AV2 85); (iii) both; (iv) *scene mix*: re-weight AV2 scenes to the nuScenes distribution of seven label-free scene factors with cross-fitted domain-classifier weights, which isolates pure covariate shift. Two further candidates were *measured* and excluded: excess history noise (a quadratic-fit residual estimate is smaller on nuScenes: 0.13 m vs. 0.29 m) and map-point density (equal per scene).

G. Label-free shift monitor

A domain classifier on label-free scene features (native sampling interval, nearby-agent count, map density, nearby lanes, history speed, heading change, curvature) is trained to separate source and target scenes; its AUC is the monitor. Within a dataset, the same classifier is trained per city against the rest. We test whether the AUC predicts the coverage loss (across cities, by Spearman correlation).

IV. Experimental Setup

a) Data: Argoverse 2 [60], nuScenes [61] and Waymo Open Motion [62] (30 of 150 validation shards), converted with ScenarioNet [63] and loaded through UniTraj [7]. AV2 scenarios are independent 11 s clips with one focal agent; the labelled AV2 validation split is divided by a fixed salted hash into a calibration half and a test half. nuScenes has a mean of 66 agents per 20 s scene (9041 from 138 validation scenes), so its calibration/test split is made by scene (64/74 scenes). Waymo scenarios carry 3.7 agents on average and are split by scenario (4331 calibration, 4318 test).

b) Forecasters: AutoBot-Ego [64] with $K = 6$ modes. The *confirmatory* models were trained after the hypotheses about them were registered, with UniTraj's recipe (effective batch 128, learning rate 7.5×10^{-4} halved at epochs 10–50): **AB-AV2-s0** on the full 183k-scene AV2 training set and **AB-NS-s0** on the full nuScenes training split. (Names read architecture–data–run: AB AutoBot, WF Wayformer; s0, s1 training seeds; r1–r3 reduced budget.) Each had to reach validation minADE_6 within 15% of UniTraj's published value [7]; their final checkpoints reach 0.846 (gate 0.98, reference 0.85) and 1.336 (gate 1.39, reference 1.21). The final model is the last checkpoint, fixed in advance. A second architecture, Wayformer [65], was trained forward-only under the same accuracy gate and effective batch size, with its own learning rate and 27-epoch schedule (reaching 0.932, Sec. V-F); no reverse-direction Wayformer was trained. Four *reduced-budget* AutoBot models (training-data subsets or shorter schedules; the three AV2 models reach minADE_6 1.09–1.35) were used to design the analyses and are labelled exploratory throughout.

c) Protocol: $\alpha = \delta = 0.1$. Forward: calibrate on the AV2 calibration half; test on the AV2 test half (same-domain control), nuScenes and Waymo. Reverse: calibrate on the nuScenes calibration scenes; test on nuScenes test scenes, AV2 and Waymo. Label-budget experiments draw k agents or m whole scenes 1000 times from a labelled pool that is disjoint from the evaluation pool. Intervals on multi-agent targets resample scenes. The source threshold is split conformal on all calibration scores: for AV2 each independent clip contributes one score, so its guarantee applies as stated; for nuScenes and Waymo the scores of agents in one scene are dependent, so the formal guarantee does not cover the source threshold itself, and we verify it empirically on a held-out same-domain split (0.4 and -1.9 points) and measure every loss against that control.

d) Registration and disclosures: Hypotheses, thresholds and success criteria were written to a version-controlled, timestamped log (released with the code) before the runs they govern; later analyses are dated addenda and outcomes are appended, never edited. Disclosed defects: a cache collision, per-agent nuScenes splits and truncated scenario identifiers (none affecting a reported number), and a betting p-value whose bet cap allowed negative capital, found in an internal review, covered by a regression test, and followed by recomputing every repair result in Sec. V-E under a design registered beforehand.

V. Results

Confirmatory results come first, organised by the four questions of Sec. II. The reduced-budget models on which the analyses were designed are marked "expl." in Table I.

A. Q1: How large is the gap?

On its own test split each confirmatory model is calibrated as promised: the same-domain gap is 0.4 points for the AV2 model and -1.9 points for the nuScenes model (Fig. 2). On every other dataset it under-covers. Calibrated on AV2, coverage falls by 15.1 points on nuScenes and by 23.4 points on Waymo. Calibrated on nuScenes, it falls by 7.7 points on AV2 and by 18.3 points on Waymo. All four intervals exclude zero by a wide margin, so a Holm correction over the four registered tests changes no conclusion. The Waymo tests were registered before any Waymo prediction was computed. In metres, the AV2-calibrated tube has radius 3.8 m, while 90% coverage needs 8.2 m on nuScenes and 8.9 m on Waymo: the source region is less than half as wide as the target requires.

The nuScenes test agents come from only 138 scenes, so they are strongly dependent: resampling whole scenes inflates the variance by a design effect of 3.9, so all intervals on multi-agent targets resample scenes. Across training budgets the more accurate model has the larger gap: the confirmatory AV2 model, within 0.5% of UniTraj's published accuracy, loses 15.1 points, against 3.3 to 10.5 points for the three reduced-budget AV2 models (Table I). The pattern does not hold between training seeds: a second seed of the same recipe is slightly less accurate (0.870) yet loses 20.6 points. It is an association across budgets of one architecture family, not a causal claim. The exploratory models of Table I, analysed before the confirmatory ones existed, reproduce the sign of every confirmatory finding with magnitudes varying by a factor of three, so the claims rest on the confirmatory models.

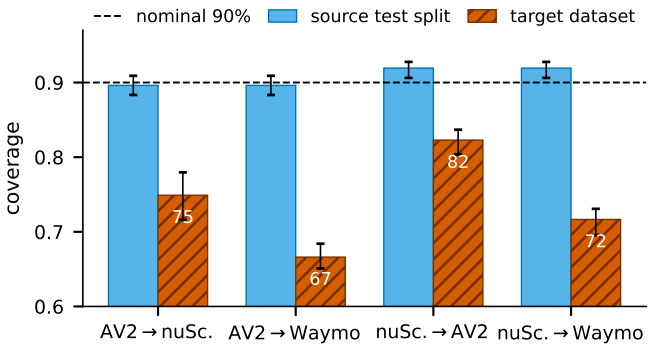


Fig. 2. Coverage of the source-calibrated region ($\alpha = 0.1$) on the source test split and on the target dataset, confirmatory models. Error bars: 95% intervals, resampled by scene where targets have several agents per scene (nuScenes, Waymo).

TABLE I
EVERY MODEL ANALYSED ($\alpha = 0.10$, POINTS; TARGET NUSCENES FOR FWD., AV2 FOR REV.). CONF.: GATED FULL-DATA MODEL; EXPL.: REDUCED-BUDGET MODEL USED TO DESIGN THE ANALYSES (NAMING: SEC. IV-B). ACC.: VALIDATION minADE_6 OF THE ANALYSED CHECKPOINT. CI: 95%, RESAMPLED BY SCENE. NORM.: GAP OF THE NORMALISED SCORE. REW.: CHANGE IN GAP FROM COVARIATE REWEIGHTING (POSITIVE = WORSE). AUDIT: REJECT PROBABILITY / FALSE-ALARM % AT $k = 1000$ LABELS.

Model	St.	Dir.	acc.	gap	95% CI	norm.	rew.	audit
AB-AV2-s0	conf.	fwd	0.846	15.1	12.0–18.4	11.0	+1.9	1.00/7
AB-AV2-s1	conf.	fwd	0.870	20.6	17.3–24.3	16.1	+1.8	1.00/23
WF-AV2-s0	conf.	fwd	0.932	6.4	3.9–9.1	1.3	+2.4	1.00/3
WF-AV2-s1	conf.	fwd	0.941	7.7	5.1–10.7	3.3	+2.5	1.00/12
AB-NS-s0	conf.	rev	1.336	7.7	6.3–9.9	8.6	+0.1	1.00/0
AB-AV2-r1	expl.	fwd	1.349	3.3	0.7–6.2	1.7	+2.0	0.95/2
AB-AV2-r2	expl.	fwd	1.347	8.0	5.0–11.3	7.0	+2.1	1.00/1
AB-AV2-r3	expl.	fwd	1.092	10.5	7.9–13.8	7.3	+2.2	1.00/6
AB-NS-r1	expl.	rev	1.300	4.7	3.6–6.2	7.1	−0.5	1.00/19

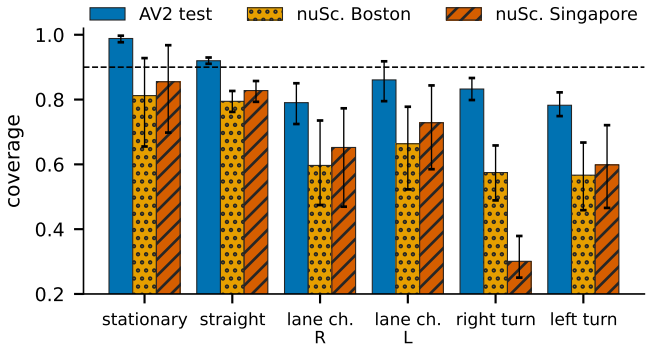


Fig. 3. Coverage by manoeuvre for the confirmatory AV2 model ($\alpha = 0.1$, dashed = nominal). nuScenes is split by city; Singapore drives on the left. lane ch. R/L: lane change to the right/left. Error bars: 95% scene-cluster bootstrap intervals (whole scenes resampled).

TABLE II
WHERE DOES THE LOSS COME FROM? COVERAGE LOST (POINTS; SHARE OF THE REAL GAP G IN PARENTHESES) WHEN ONE MEASURED AV2–NUSCENES DIFFERENCE IS INJECTED INTO AV2 TEST INPUTS: 2 HZ HISTORY, LANE-GRAPH DROPOUT, BOTH, AND RE-WEIGHTING AV2 SCENES TO NUSCENES' SCENE MIX. EXCESS ANNOTATION NOISE AND MAP-POINT DENSITY WERE MEASURED AND FOUND ABSENT (NUSCENES IS SMOOTHER; EQUAL DENSITY), SO THEY ARE NOT INJECTED.

Model	Set	G (pt)	2 Hz	lanes	both	scene mix
AB-AV2-r1	expl.	+3.3	+0.6 (19%)	+1.4 (43%)	+1.8 (53%)	−1.3 (−38%)
AB-AV2-r2	expl.	+8.0	+0.7 (8%)	+1.4 (17%)	+1.4 (17%)	−1.4 (−18%)
AB-AV2-r3	expl.	+10.5	+1.1 (11%)	+1.3 (12%)	+2.2 (21%)	−1.1 (−11%)
AB-AV2-s0	conf.	+15.1	+2.3 (15%)	+1.8 (12%)	+4.1 (27%)	−1.0 (−6%)

B. Q2: Does covariate reweighting fix it?

No. Weighting the calibration scores by an estimated, label-free likelihood ratio increased the gap on every forward model by about two points (column "rew." in Table I) and left the reverse direction unchanged. Reweighting AV2 scenes to the nuScenes mix of speeds and manoeuvres even moves coverage the other way (Table II): nuScenes' scene mix is easier, so composition cannot explain the loss. This is the pattern expected when the conditional distribution of the future given the scene changes, which no reweighting on the scene can repair (Proposition 1).

C. Where does the loss come from?

The loss is concentrated in turns. On the confirmatory model, right-turn coverage falls by 41 points from AV2 to nuScenes, against 11 points for straight driving (Fig. 3). It is largest in Singapore, where traffic drives on the left: the right-turn drop there exceeds the drop in Boston by 27 points. What AV2-trained forecasters miss is the wide, across-traffic turn of left-hand traffic. For AutoBot, injecting one measured input difference at a time into AV2 test inputs reproduces only part of the gap (Table II): a 2 Hz history reproduces 15%, lane dropout 12%, and both together about a quarter. Excess annotation noise and map-point density were measured and found absent.

D. Q3: What can label-free methods do?

Little, and not reliably. The model-normalised score (Sec. III-A) reduces the AV2→nuScenes gap by 27% for AutoBot (to 11.0 points) and by 80% for Wayformer, but widens the reverse gap (to 8.6 points). On the reduced-budget models its effect ranges from a 49% reduction to a 52% increase. A domain-classifier shift monitor correlates with per-city coverage loss at Spearman $\rho \leq 0.47$ on every model, below the registered bar of 0.5. Adaptive conformal inference [13] converges only with a stream of labels; frozen after 1000 labels, its threshold covers $\geq 90\%$ of held-out scenes in only 40–50% of streams. These results agree with Proposition 1: without target labels there is no general fix.

One label-free method restricts the shift class, as that proposition says a useful one must. *Protocol-matched recalibration* passes the source calibration inputs through the measured AV2→nuScenes input differences of Sec. III-F (2 Hz history and lane dropout, estimated from target inputs only) and recalibrates on the result. It is exploratory and has no guarantee. For AutoBot it moves the threshold from 3.8 to 4.7 m and coverage on nuScenes from 75% to 81%, closing 41% of the gap. For Wayformer it moves it from 4.2 to 5.7 m, coverage from 84% to 92%, and overshoots the gap to −2.3 points, in the conservative direction, which may partly reflect how Wayformer's velocity inputs are reconstructed for injected trajectories (Sec. V-F). It does not shorten the label budget of Sec. V-E: started from the matched threshold, the scene-level rule still needs $m = 150$ scenes (AutoBot) and $m = 150$ (Wayformer) for a usable region. A measured correction recovers part of the gap but does not replace labels.

The loss is not an artefact of the registered score. Re-scored with four alternatives (exploratory: top-probability mode, final displacement, probability-weighted, and per-time-step scaling by the source median error), all 24 out-of-domain cases undercover, by 4 to 26 points. No score is uniformly better: the time-scaled score cuts AutoBot's AV2→Waymo loss from 23 to 9 points but raises its AV2→nuScenes loss from 15 to 24.

Further label-free constructions, registered before running, also fail. *Worst-domain calibration* (the largest threshold over the other labelled datasets) meets its bar (gap ≤ 2 points at $\leq 2\times$ the oracle area) in only 2 of six held-out cells, over-covering easier targets by up to 8 points and under-covering Waymo, the hardest dataset, by 3–15. *Manoeuvre-conditional calibration*, with groups from history only, closes a median of 1% of the gap (bar met in 0 of six). Robust conformal prediction [28] with a leave-one-dataset-out radius meets the bar in 2 of six, over-covering by up to 8 points and under-covering Waymo by up to 10. A restricted shift class helps only if the target belongs to it (Proposition 1).

One label-free signal does track the loss: the drift of the forecaster's own predicted scale (log ratio S of its median predicted Laplace scale on the target to that on the source calibration split) correlates with the coverage gap at Spearman $\rho = 0.77$ over 24 model–target pairs (0.80 over the 14 shifted ones), and the Kolmogorov–Smirnov distance of the two scale distributions at 0.92, against ≤ 0.47 for the domain classifier. The pairs are not independent (shared models and targets), within nuScenes alone the KS distance ranks the models at $\rho = 0.90$ (post hoc), and it is blind to the nuScenes→AV2 shift, which loses 8.1 points while the KS distance (0.04) stays within the same-domain range (up to 0.05). It is a warning

signal, not a guarantee.

Training on more data helps but does not close the gap. **AB-AV2+NS-s0**, trained on AV2 and nuScenes and calibrated on both, loses 15.2 points (interval 14.5–15.9) on Waymo, which it never saw, against 23.4 for AB-AV2-s0 (-35%) and 18.3 for AB-NS-s0, with in-domain gaps of 2.1 and 1.5 points; it misses our registered bar of 11.7 points. Its nuScenes accuracy improves (1.094 against 1.336) and AV2 accuracy is 0.890. The largest per-dataset threshold would give 11.0 points on Waymo but over-covers AV2 by 4.3, failing the in-domain condition.

E. Q4: How many labels to audit and to repair?

We report three label budgets: to *detect* the failure, to *certify* a repaired threshold, and for that threshold to be *useful*. The second does not imply the third.

Audit (detect). With $k = 1000$ labelled target agents (above the $\approx$ 750 that Eq. (3) requires for a 3.4-point shortfall, close to the smallest gap in Table I) the exact binomial audit (Sec. III-B) detects the gap with probability 1.00 on every confirmatory run (Table I). On a shift-free split it rejects in 0–23% of label draws, because it tests the realised source threshold, which under-covers its own test split by up to 1.0 points when its calibration draw is unlucky; a deployment that only wants to flag shifts should test $1-\alpha-\epsilon$ for a small margin ϵ.

Repair with exchangeable labels (certify). When $k = 1000$ agents are labelled at random, certify-or-recalibrate (Sec. III-C) covers $\geq 90\%$ of held-out agents in 94–98% of label draws, against 50–60% for plain recalibration, at a region 1.18–1.23$\times$ the oracle's.

Repair with whole-scene labels (certify at the scene level). Real labels arrive scene by scene, which breaks agent-level exchangeability. We draw m labelled scenes from a pool disjoint from the evaluation scenes (nuScenes train for AV2→nuScenes, the Waymo calibration half for the Waymo targets) and evaluate the scene-averaged miscoverage $R(\hat{q})$ exactly on the evaluation pool, 1000 draws per m (Fig. 4). The scene-level rule with the betting p-value is valid at every m on all four transfers ($\Pr[R(\hat{q}) \leq \alpha] \geq 0.94$). The agent-level rule is not: on AV2→nuScenes, with a mean of 66 agents per scene (median 53), it is valid in 56% of draws at $m = 40$ and 4% at $m = 300$, and one-agent-per-scene split CP [35] in 57%. With few agents per scene it stays valid in at least 94% of draws for AV2→Waymo, nuScenes→AV2 and Wayformer's AV2→Waymo, but not for nuScenes→Waymo, where it falls from 88% at $m = 40$ to 65% at $m = 300$ although Waymo scenes are small: there the source threshold misses 27.8% of evaluation agents. Agent-level guarantees fail where scenes are large or the source threshold is badly wrong (Sec. V-F).

Useful. Validity is not usefulness. With $m \leq 40$ scenes the scene-level rule abstains in every draw, correctly reporting that it cannot certify any finite threshold rather than returning an unsupported one. On AV2→nuScenes it abstains in 57% of draws at $m = 80$ and still 6% at $m = 100$. It returns a usable region (at most twice the oracle area, abstaining in $\leq 5\%$ of draws) from $m = 150$ scenes for AV2→nuScenes,

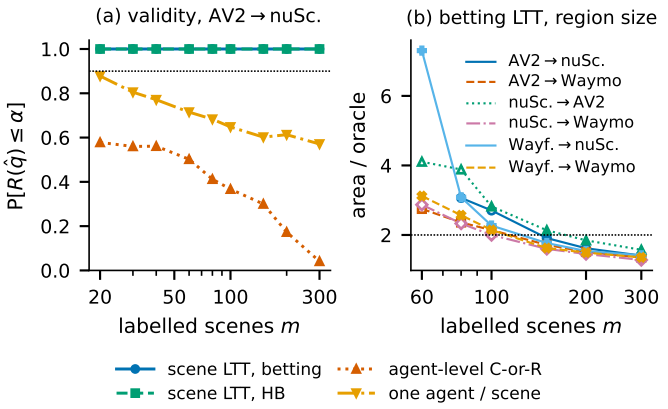


Fig. 4. Repair with m whole labelled scenes, disjoint label and evaluation pools, 1000 draws per m. (a) Probability that the repaired threshold controls scene-averaged miscoverage at α (dotted line: $1-\delta$), AutoBot only. (b) Median region area relative to the oracle (area proxy, Sec. III-A), all six transfers of both architectures; hollow markers: threshold infinite in $> 5\%$ of draws.

AV2→Waymo and nuScenes→Waymo, and from $m = 200$ for nuScenes→AV2, whose single-agent scenes carry less information each. At that budget the repaired AV2→nuScenes radius is 11.3 m against an oracle of 8.2 m. The betting p-value gives smaller regions than Hoeffding–Bentkus at large m (1.6$\times$ vs. 2.1$\times$ the oracle area at $m = 200$ on AV2→nuScenes). In dataset terms, 150 scenes is more than the 138 scenes of the entire nuScenes validation split, so on that target the repair draws on training-split labels, which assumes nuScenes training and validation scenes are exchangeable. A deployment that can label fewer scenes should expect the rule to abstain, not to return a smaller guaranteed region. Stopping early (the same guarantee holds at any stopping time) cuts the median budget, once the threshold is stable, to 66–106 scenes (33–71% of the fixed budget, validity at least 0.996), but the regions are then 2.1–4.7$\times$ the oracle area, above our 2$\times$ bar.

F. A second architecture: Wayformer

To test whether the findings are AutoBot-specific we trained Wayformer [65] on the same full AV2 set under the same gate (val minADE_6 0.932 for the final checkpoint; UniTraj reports 0.85 for both architectures [7, Table 15]), forward only.

The direction replicates everywhere: same-domain gap 0.1 points; AV2→nuScenes gap $+6.4$ points (CI [3.9, 9.1]); AV2→Waymo gap $+25.0$ points (CI [23.5, 26.9]). Reweighting again makes the nuScenes gap worse (+2.4 points); the audit has power 1.00. The loss again concentrates in turns (right-turn drop 20 points vs. 5.5 for straight) and in Singapore (14 points more than Boston), with a similar scene-cluster design effect (3.0 vs. AutoBot's 3.9). The scene-level repair is valid at every m on both Wayformer transfers (Fig. 4b) and reaches the same operating point as AutoBot: $m = 150$ scenes for AV2→nuScenes (area 1.8$\times$) and $m = 150$ for AV2→Waymo (area 1.6$\times$).

The input-level explanation does not replicate in size. Injecting the 2 Hz history costs Wayformer 8.2 points (129% of its nuScenes gap), and with lane dropout 10.6 points (166%),

against about a quarter for AutoBot. Wayformer reads velocity, acceleration and heading channels that AutoBot ignores and that we reconstruct by finite difference for an injected trajectory, which may overstate the effect; lane dropout, which needs no reconstruction, still gives 16%.

Magnitude does not replicate. Wayformer's AV2→nuScenes gap is under half of AutoBot's 15.1 points, while its AV2→Waymo gap (+25.0) is close to AutoBot's (+23.4), so the "more accurate model, larger gap" pattern (Sec. V-A) breaks: Wayformer is about as accurate yet has a far smaller nuScenes gap. A second, unregistered difference: on the *same* nuScenes target, agent-level certify-or-recalibrate was almost useless for AutoBot (56% valid at $m = 40$, 4% at $m = 300$) but mostly valid for Wayformer (82%, 96%) on identical scenes. It tracks the size of the miscalibration: the source threshold misses 26.8% of AutoBot's evaluation agents against 17.5% of Wayformer's, so scene clustering has more room to hurt a method that was already more wrong. Scene size alone does not predict agent-level failure.

VI. Discussion

A planner that trusts a source-calibrated region on a new dataset trusts a region that is too small. At $\alpha = 0.1$ the confirmatory models lose 15.1, 7.7, 23.4 and 18.3 points of coverage on the four transfers we tested. Without labels the loss cannot be measured, although the drift of the forecaster's predicted scale warns of it (Sec. V-D). Four results bear directly on practice.

Label-free fixes are not safe defaults. Reweighting made every forward gap larger, and the normalised score helped in one direction and hurt in the other.

Auditing is cheap; repairing is not. About 10^3 labelled agents detect a gap of this size; repairing it with a guarantee that survives scene clustering needs roughly 150 to 200 whole scenes before the region is finite and within twice the oracle area, and below that budget the right behaviour is to abstain (Fig. 5).

A learned predictor of scene loss does not replace labels. A source-trained classifier of whether an agent is missed explained at most 16% of the variance of the per-scene miss fraction on any target.

The unit of labelling matters, and scene size alone does not predict the failure. On nuScenes (mean 66 agents per scene), agent-level guarantees were nearly useless for AutoBot (56% valid at $m = 40$) but mostly valid for Wayformer on the identical scenes (82%, Sec. V-F), tracking how badly each model's source threshold was miscalibrated there (26.8% vs. 17.5% missed), not the scene structure, which was the same for both. Agent-level guarantees held on AV2 and on Waymo where scenes are small and the source threshold less wrong, but not for nuScenes→Waymo (65% valid at $m = 300$), whose source threshold misses 27.8% of agents. Without labels a practitioner cannot tell which case applies, so the scene-level guarantee should be the default.

a) What a planner would see: Descriptive, agent-weighted numbers (no guarantee covers them): at the source threshold, 25% of nuScenes agents and 34% of Waymo agents fall outside the AV2-calibrated region, against a nominal 10%, and the shortfall is concentrated in turns: 70% of right-turning agents in Singapore and 43% in Boston, against 21% of straight-driving agents in Boston. The tube is 3.8 m wide where 8.2 m is needed; scene-level repair restores 11.3 m.

1) *Audit.* Binomial test of coverage $\geq 1 - \alpha - \epsilon$ on $k \approx 10^3$ agents (Sec. III-B). If not rejected, keep $\hat{q}_{\mathrm{src}}$ and keep auditing.
2) *Certify.* Test $H : R(\hat{q}_{\mathrm{src}}) > \alpha$ on the m scenes at level $\delta/2$ with the betting p-value; if rejected, keep $\hat{q}_{\mathrm{src}}$.
3) *Recalibrate.* Otherwise run scene-level Learn-then-Test on the fixed grid (Sec. III-D) at level $\delta/2$ and return its threshold.
4) *Abstain.* If no finite threshold is certified, return $+\infty$ and label more scenes (expect this below about 150 scenes).

Fig. 5. Audit-then-repair protocol for a transferred conformal region. Steps 2–4 give $\Pr[R(\hat{q}) \leq \alpha] \geq 1 - \delta$ (Proposition 3); step 1 only detects.

b) Mechanism: For AutoBot the measured input differences reproduce about a quarter of the gap, and the loss concentrates in turns under left-hand traffic, which AV2 does not contain; we read this as a change in the conditional behaviour of road users rather than in the scenes (an interpretation, consistent with Proposition 1). The turn pattern replicates on Wayformer but the input-level explanation does not: its injected differences reproduce the whole, smaller gap (Sec. V-F), so how much of the loss is protocol and how much behaviour depends on the model. The accuracy–gap association across AutoBot training budgets breaks between seeds and across architectures, so accuracy is not a general predictor of the gap.

c) Limitations: (i) Two architectures; Wayformer was only trained forward and its injection result rests on reconstructed kinematics. (ii) For AutoBot the measured input differences leave most of the gap unexplained. (iii) Coverage is marginal over agents or scenes, not conditional on a scenario; the manoeuvre breakdown is a diagnostic. (iv) Proposition 1 covers scale shifts only. (v) The repair budget is measured on three datasets and may differ on others; reuse the registered design, not the numbers. (vi) Seeds: a second AutoBot seed gave an AV2→nuScenes loss of 20.6 points (95% interval 17–24) against 15.1 (12–18), meeting our registered criterion (same sign, within a factor of two); seed-to-seed variation is as large as one run's sampling uncertainty. A second Wayformer seed (accuracy 0.941) gave 7.7 points on nuScenes against 6.4 and 24.2 against 25.0 on Waymo, so its smaller nuScenes gap is not a seed artefact. The reverse direction, the injections, recalibration and the multi-dataset model were not re-seeded.

VII. Conclusion

A conformal guarantee for trajectory prediction does not transfer across driving datasets. On three datasets, two architectures and six model–target transfers, a forecaster calibrated at 90% covered between 65 and 84% of target futures, every time under-covering, never over. Reweighting does not repair the loss and no label-free method can in general, though the drift of the forecaster's predicted scale warns of it, and training on two datasets cuts the loss on a third by about a third

without removing it. More accurate AutoBot models lost more coverage across training budgets, but not between seeds or across architectures, so accuracy is not a general predictor of how badly the guarantee fails. Labelled target data can repair it: a scene-level Learn-then-Test rule with a correctly capped betting p-value restores the guarantee at every budget we tested, on both architectures, and gives a usable region from about 150 labelled scenes. We recommend that any transferred conformal region be audited on target labels and, where it fails, recalibrated at the scene level (Fig. 5) before it is used for planning.

Data and Code Availability

Argoverse 2, nuScenes and Waymo Open Motion are third-party datasets available under their own licences from https://www.argoverse.org, https://www.nuscenes.org and https://waymo.com/open. Code, configuration files, the registered hypotheses, the experiment log and all aggregate result files (no raw third-party data) are at https://github.com/hwaqas160/conformal-coverage-transfer and archived under doi:10.5281/zenodo.23054412 (CC BY 4.0). Every reported result, table and figure is regenerated from the result files by the `paper/make_results*.py` scripts; dataset descriptions (AV2 training-set size, the 30 of 150 Waymo shards, scene length) come from the dataset documentation and our conversion logs.

Conflict of Interest

The authors declare no competing financial interests or personal relationships that could have influenced this work.

Acknowledgment

Generative AI was used: the corresponding author used Anthropic's Claude (Claude Code; Sonnet and Opus models). (i) *Text:* Claude drafted and revised all sections (abstract to conclusion); the authors reviewed, edited and take full responsibility for every statement. (ii) *Code and experiments:* Claude wrote and debugged the analysis code and the scripts that generate every table, figure and number from computed results, and ran and monitored the experiments; the authors set the research question and scope and directed the work. (iii) *Figures:* generated by that code; no AI image model was used.

References

[1] V. Vovk, A. Gammerman, and G. Shafer, *Algorithmic Learning in a Random World*. Springer, 2005.

[2] A. N. Angelopoulos and S. Bates, "Conformal prediction: A gentle introduction," *Foundations and Trends in Machine Learning*, vol. 16, no. 4, pp. 494–591, 2023.

[3] L. Lindemann, M. Cleaveland, G. Shim, and G. J. Pappas, "Safe planning in dynamic environments using conformal prediction," *IEEE Robotics and Automation Letters*, vol. 8, no. 8, pp. 5116–5123, 2023.

[4] S. H. Sun and R. Yu, "Copula conformal prediction for multi-step time series prediction," in *International Conference on Learning Representations (ICLR)*, 2024.

[5] H. Huang, S. He, and F. Miao, "CUQDS: Conformal uncertainty quantification under distribution shift for trajectory prediction," in *Proceedings of the AAAI Conference on Artificial Intelligence*, 2025.

[6] L. Lindemann, Y. Zhao, X. Yu, G. J. Pappas, and J. V. Deshmukh, "Formal verification and control with conformal prediction," arXiv:2409.00536, 2024.

[7] L. Feng, M. Bahari, K. M. B. Amor, É. Zablocki, M. Cord, and A. Alahi, "UniTraj: A unified framework for scalable vehicle trajectory prediction," in *European Conference on Computer Vision (ECCV)*, 2024.

[8] T. Demmler, J. Häringer, A. Tamke, T. Dang, A. Hegai, and L. Mikelsons, "Beyond features: How dataset design influences multi-agent trajectory prediction performance," arXiv:2507.05098, 2025.

[9] Y. Ovadia, E. Fertig, J. Ren, Z. Nado, D. Sculley, S. Nowozin, J. V. Dillon, B. Lakshminarayanan, and J. Snoek, "Can you trust your model's uncertainty? Evaluating predictive uncertainty under dataset shift," in *Advances in Neural Information Processing Systems (NeurIPS)*, 2019.

[10] R. J. Tibshirani, R. F. Barber, E. J. Candès, and A. Ramdas, "Conformal prediction under covariate shift," in *Advances in Neural Information Processing Systems (NeurIPS)*, 2019.

[11] R. F. Barber, E. J. Candès, A. Ramdas, and R. J. Tibshirani, "Conformal prediction beyond exchangeability," *The Annals of Statistics*, vol. 51, no. 2, pp. 816–845, 2023.

[12] L. Aolaritei, Z. O. Wang, J. Zhu, M. I. Jordan, and Y. Marzouk, "Conformal prediction under Lévy–Prokhorov distribution shifts: Robustness to local and global perturbations," arXiv:2502.14105, 2025.

[13] I. Gibbs and E. J. Candès, "Adaptive conformal inference under distribution shift," in *Advances in Neural Information Processing Systems (NeurIPS)*, 2021.

[14] A. N. Angelopoulos, E. J. Candès, and R. J. Tibshirani, "Conformal PID control for time series prediction," in *Advances in Neural Information Processing Systems (NeurIPS)*, 2023.

[15] Y. Shu, J. Xu, L. Kong, F. Zhang, G. Yin, and C. Sun, "Scenario-aware uncertainty quantification for trajectory prediction with statistical guarantees," arXiv:2512.05682, 2025.

[16] X. Chen, R. Bhadani, and L. Head, "Conformal trajectory prediction with multi-view data integration in cooperative driving," arXiv:2408.00374, 2024.

[17] N. A. Madjid, A. Ahmad, M. Mebrahtu, Y. Babaa, A. Nasser, S. Malik, B. Hassan, N. Werghi, J. Dias, and M. Khonji, "Trajectory prediction for autonomous driving: Progress, limitations, and future directions," arXiv:2503.03262, 2025.

[18] F. C. Pouria, M. Golchoubian, and K. Driggs-Campbell, "Rethinking Gaussian trajectory predictors: Calibrated uncertainty for safe planning," arXiv:2603.10407, 2026.

[19] J. Lei, M. G'Sell, A. Rinaldo, R. J. Tibshirani, and L. Wasserman, "Distribution-free predictive inference for regression," *Journal of the American Statistical Association*, vol. 113, no. 523, pp. 1094–1111, 2018.

[20] Z. Mei, A. Dixit, M. Booker, E. Zhou, M. Storey-Matsutani, A. Z. Ren, O. Shorinwa, and A. Majumdar, "Perceive with confidence: Statistical safety assurances for navigation with learning-based perception," arXiv:2403.08185, 2024.

[21] J. Yao, X. Zhang, Y. Xia, Z. Wang, A. K. Roy-Chowdhury, and J. Li, "Towards generalizable safety in crowd navigation via conformal uncertainty handling," in *Conference on Robot Learning (CoRL)*, 2025.

[22] A. O. Feldman, J. A. Vincent, M. Adang, J. Low, and M. Schwager, "Learning robot safety from sparse human feedback using conformal prediction," arXiv:2501.04823, 2025.

[23] A. E. Binny, M. Anand, H. T. M. Kussaba, L. Chen, S. Agrawal, F. J. Abu-Dakka, and A. Swikir, "Safe and stable neural network dynamical systems for robot motion planning," *IEEE Robotics and Automation Letters*, 2026.

[24] O. Mirzaeedodangeh, E. Shekhtman, N. Matni, and L. Lindemann, "Safe planning in interactive environments via iterative policy updates and adversarially robust conformal prediction," arXiv:2511.10586, 2025.

[25] A. E. Binny and A. Dixit, "Who moved my distribution? Conformal prediction for interactive multi-agent systems," arXiv:2511.11567, 2025.

[26] R. Xu, C. Chen, Y. Sun, P. Venkitasubramaniam, and S. Xie, "Wasserstein-regularized conformal prediction under general distribution shift," arXiv:2501.13430, 2025.

[27] A. H. C. Correia and C. Louizos, "Non-exchangeable conformal prediction with optimal transport: Tackling distribution shifts with unlabeled data," in *Advances in Neural Information Processing Systems (NeurIPS)*, 2025.

[28] M. Cauchois, S. Gupta, A. Ali, and J. C. Duchi, "Robust validation: Confident predictions even when distributions shift," *Journal of the American Statistical Association*, vol. 119, no. 548, pp. 3033–3044, 2024.

[29] J. Liu, D. Qiao, Y. Zhang, and X. Chen, "Certify or refuse: A cross-model map for selective risk control with coverage floors under covariate shift," arXiv:2608.10893, 2026.
[30] R. Tumu, A. Singh, and R. Mangharam, "AdaptNC: Adaptive non-conformity scores for conformal prediction under distribution shift," arXiv:2602.01629, 2026.
[31] J. Jun and I. Ohn, "Online conformal inference with retrospective adjustment for faster adaptation to distribution shift," *Pattern Recognition*, vol. 180, 2026.
[32] T. Liu, E. Dobriban, and F. Orabona, "Online conformal prediction via universal portfolio algorithms," arXiv:2602.03168, 2026.
[33] J. Liang, Z. Ren, and Y. Chen, "Optimal training-conditional regret for online conformal prediction," arXiv:2602.16537, 2026.
[34] K. Rahaman, J. V. Deshmukh, A. R. Hota, and L. Lindemann, "When environments shift: Safe planning with generative priors and robust conformal prediction," arXiv:2602.12616, 2026.
[35] R. Dunn, L. Wasserman, and A. Ramdas, "Distribution-free prediction sets for two-layer hierarchical models," *Journal of the American Statistical Association*, vol. 118, no. 544, pp. 2491–2502, 2023.
[36] Y. Lee, R. F. Barber, and R. Willett, "Distribution-free inference with hierarchical data," arXiv:2306.06342, 2023.
[37] G. Principato, G. Stoltz, Y. Amara-Ouali, Y. Goude, B. Hamrouche, and J.-M. Poggi, "Conformal prediction for hierarchical data," arXiv:2411.13479, 2024.
[38] A. Bhattacharyya and R. F. Barber, "Group-weighted conformal prediction," arXiv:2401.17452, 2024.
[39] S. Mallick, E. T. Tchetgen, E. Dobriban, and Y. Lee, "Generalized hierarchical conformal prediction," arXiv:2608.15500, 2026.
[40] Y. Min, L. Peng, and C. Zou, "A unified theory of conditional coverage in conformal prediction with applications," arXiv:2605.11602, 2026.
[41] Z. Zhou, X. Zhang, C. Tao, and Y. Yang, "Conformal prediction assessment: A framework for conditional coverage evaluation and selection," arXiv:2603.27189, 2026.
[42] S. Braun, D. Holzmüller, M. I. Jordan, and F. Bach, "Conditional coverage diagnostics for conformal prediction," arXiv:2512.11779, 2025.
[43] K. Bairaktari, J. Wu, and Z. S. Wu, "Kandinsky conformal prediction: Beyond class- and covariate-conditional coverage," arXiv:2502.17264, 2025.
[44] S. Bates, A. Angelopoulos, L. Lei, J. Malik, and M. Jordan, "Distribution-free, risk-controlling prediction sets," *Journal of the ACM*, vol. 68, no. 6, pp. 43:1–43:34, 2021.
[45] A. N. Angelopoulos, S. Bates, E. J. Candès, M. I. Jordan, and L. Lei, "Learn then test: Calibrating predictive algorithms to achieve risk control," *The Annals of Applied Statistics*, vol. 19, no. 2, 2025.
[46] I. Waudby-Smith and A. Ramdas, "Estimating means of bounded random variables by betting," *Journal of the Royal Statistical Society: Series B (Statistical Methodology)*, vol. 86, no. 1, pp. 1–27, 2024.
[47] J. Alvarez, "A distribution-free valid p-value for finite samples of bounded random variables," arXiv:2405.08975, 2024.
[48] D. T. Nguyen, R. Pathak, A. N. Angelopoulos, S. Bates, and M. I. Jordan, "Data-adaptive tradeoffs among multiple risks in distribution-free prediction," arXiv:2403.19605, 2024.
[49] B.-S. Einbinder, L. Ringel, and Y. Romano, "Semi-supervised risk control via prediction-powered inference," arXiv:2412.11174, 2024.
[50] S. Joshi, T. Wang, H. Hassani, and E. Dobriban, "Risk-controlled post-processing of decision policies," arXiv:2605.06479, 2026.
[51] E. Gauthier, F. Bach, and M. I. Jordan, "e-values expand the scope of conformal prediction," arXiv:2503.13050, 2025.
[52] D. Prinster, X. Han, A. Liu, and S. Saria, "WATCH: Adaptive monitoring for AI deployments via weighted-conformal martingales," in *International Conference on Machine Learning (ICML)*, 2025.
[53] Q. Wu and M. Campbell, "Semantic and feature guided uncertainty quantification of visual localization for autonomous vehicles," in *IEEE International Conference on Robotics and Automation (ICRA)*, 2025.
[54] N. Wang, D. Shang, Y. Gong, X. Hu, Z. Song, L. Yang, Y. Huang, X. Wang, and J. Lu, "Collaborative perception datasets for autonomous driving: A review," arXiv:2504.12696, 2025.
[55] L. Kong, X. Xu, Y. Liu, J. Cen, R. Chen, W. Zhang, L. Pan, K. Chen, and Z. Liu, "LargeAD: Large-scale cross-sensor data pretraining for autonomous driving," *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 2025.
[56] X. Qiu, W. Liu, F. Ai, Y. Tan, Z. Xu, and C. Song, "Understanding cross-sensor feature variations for generalizable 3D perception," arXiv:2606.11573, 2026.
[57] V. Vovk, "Conditional validity of inductive conformal predictors," in *Proceedings of the Asian Conference on Machine Learning (ACML)*, ser. Proceedings of Machine Learning Research, vol. 25, 2012, pp. 475–490.
[58] M. Bian and R. F. Barber, "Training-conditional coverage for distribution-free predictive inference," *Electronic Journal of Statistics*, vol. 17, no. 2, pp. 2044–2066, 2023.
[59] C. J. Clopper and E. S. Pearson, "The use of confidence or fiducial limits illustrated in the case of the binomial," *Biometrika*, vol. 26, no. 4, pp. 404–413, 1934.
[60] B. Wilson, W. Qi, T. Agarwal, J. Lambert, J. Singh, S. Khandelwal, B. Pan, R. Kumar, A. Hartnett, J. K. Pontes, D. Ramanan, P. Carr, and J. Hays, "Argoverse 2: Next generation datasets for self-driving perception and forecasting," in *NeurIPS Datasets and Benchmarks Track*, 2021.
[61] H. Caesar, V. Bankiti, A. H. Lang, S. Vora, V. E. Liong, Q. Xu, A. Krishnan, Y. Pan, G. Baldan, and O. Beijbom, "nuScenes: A multimodal dataset for autonomous driving," in *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2020.
[62] S. Ettinger, S. Cheng, B. Caine, C. Liu, H. Zhao, S. Pradhan, Y. Chai, B. Sapp, C. R. Qi, Y. Zhou, Z. Yang, A. Chouard, P. Sun, J. Ngiam, V. Vasudevan, A. McCauley, J. Shlens, and D. Anguelov, "Large scale interactive motion forecasting for autonomous driving: The Waymo open motion dataset," in *IEEE/CVF International Conference on Computer Vision (ICCV)*, 2021.
[63] Q. Li, Z. Peng, L. Feng, Z. Liu, C. Duan, W. Mo, and B. Zhou, "ScenarioNet: Open-source platform for large-scale traffic scenario simulation and modeling," in *NeurIPS Datasets and Benchmarks Track*, 2023.
[64] R. Girgis, F. Golemo, F. Codevilla, M. Weiss, J. A. D'Souza, S. E. Kahou, F. Heide, and C. Pal, "Latent variable sequential set transformers for joint multi-agent motion prediction," in *International Conference on Learning Representations (ICLR)*, 2022.
[65] N. Nayakanti, R. Al-Rfou, A. Zhou, K. Goel, K. S. Refaat, and B. Sapp, "Wayformer: Motion forecasting via simple & efficient attention networks," in *IEEE International Conference on Robotics and Automation (ICRA)*, 2023, pp. 2980–2987.

Hassan Waqas Saleem received his B.Tech. (Pass) in Electrical Technology from Bahauddin Zakariya University, Multan, Pakistan, in 2014, his B.Sc. in Electrical Engineering from the University of Engineering and Technology, Taxila, Pakistan, in 2018, and his M.S. in Electrical Engineering (Artificial Intelligence and Autonomous Systems) from the School of Electrical Engineering and Computer Science (SEECS), National University of Sciences and Technology (NUST), Islamabad, Pakistan, in 2024. He currently serves as Manager (Technical) at a public sector organization in Pakistan. His research topics mainly include artificial intelligence, graph neural networks, large language models, knowledge graph completion, robotics, and simultaneous localization and mapping (SLAM).

Rahat Zeeshan Saleem received his MBA from Bahauddin Zakariya University, Multan, Pakistan, in 2021, graduating as a Gold Medalist, and his M.S. in Supply Chain Management from NUST Business School (NBS), National University of Sciences and Technology (NUST), Islamabad, Pakistan, in 2023. He currently serves as Project Lead of the Supply Chain Management team at a private sector organization in Pakistan. His research topics mainly include supply chain management, sustainable and resilient supply networks, supplier selection and substitution, and applications of artificial intelligence to supply chain decision support.

Hina Maryam received her M.Phil. in Statistics from Bahauddin Zakariya University, Multan, Pakistan, in 2019, graduating as a Gold Medalist. She currently serves as Lecturer at Kinnaird College for Women University, Lahore, Pakistan. Her research topics mainly include statistical modelling, data optimization, applied data analysis, and quantitative methods for economic and financial data.